



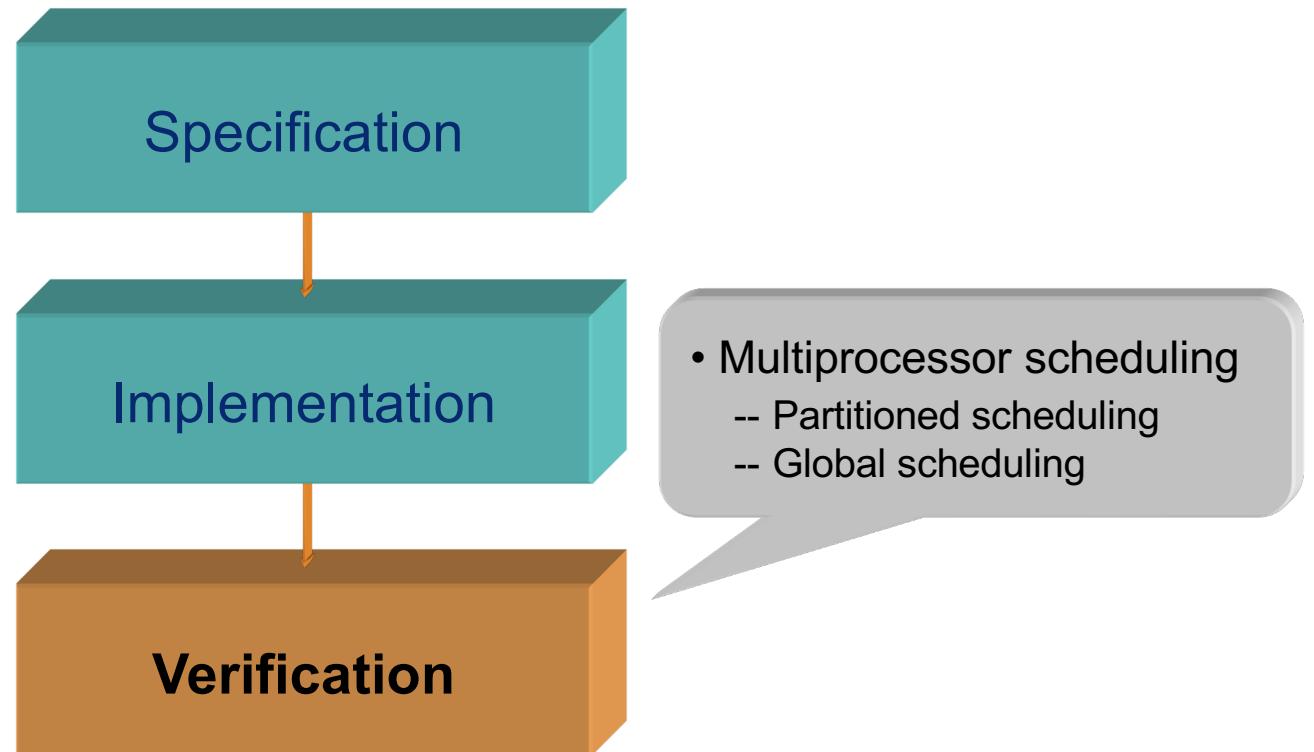
# Real-Time Systems

Lecture #14

Docent Risat Pathan

Department of Computer Science and Engineering  
Chalmers University of Technology

# Real-Time Systems



# Multiprocessor scheduling

## How are tasks assigned to processors?

- Static assignment
  - The processor(s) used for executing a task are determined before system is put in mission (“off-line”)
  - Approach: **Partitioned scheduling**
- Dynamic assignment
  - The processor(s) used for executing a task are determined during system operation “on-line”
  - Approach: **Global scheduling**

# Multiprocessor scheduling

## How are tasks allowed to migrate?

- Partitioned scheduling
  - No migration!
  - Each instance of a task must execute on the same processor
  - Equivalent to multiple single-processor systems!
- Global scheduling
  - Full migration!
  - A task is allowed to execute on an arbitrary processor
  - Migration can occur even during execution of an instance of a task (for example, after being preempted)

# Multiprocessor scheduling

A fundamental limit: (Andersson, Baruah & Jonsson, 2001)

The utilization guarantee bound for multiprocessor scheduling (partitioned or global), using task priorities only, cannot be higher than 50% of the capacity of the processors.

- Hence, we should not expect to utilize more than half the processing capacity if hard real-time constraints exist.
- A way to circumvent this limit is to use p-fair (priorities + time quanta) scheduling and dynamic task priorities.

# Partitioned scheduling

## General characteristics:

- Each processor has its own queue for ready tasks
- Tasks are organized in groups, and each task group is assigned to a specific processor
  - For example, using a bin-packing algorithm
- When selected for execution, a task can only be dispatched to its assigned processor

# Partitioned scheduling

## Advantages:

- Mature scheduling framework
  - Most single-processor scheduling theory also applicable here
  - Single-processor resource-management protocols can be used
- Supported by automotive industry
  - AUTOSAR prescribes partitioned scheduling

## Disadvantages:

- Cannot exploit all unused execution time
  - Surplus capacity cannot be shared among processors
  - Will suffer from overly-pessimistic WCET derivation

# Partitioned scheduling

## Bin-packing algorithms:

- Basic idea:
  - The problem concerns packing objects of varying sizes in boxes ("bins") with the objective of minimizing number of used boxes.
- Application to multiprocessor systems:
  - Bins are represented by processors and objects by tasks.
  - The decision whether a processor is "full" or not is derived from a utilization-based feasibility test.
- Assumptions:
  - Independent, periodic tasks
  - Preemptive, single-processor scheduling (RM)



# Partitioned scheduling

## Bin-packing algorithms:

Rate-Monotonic-First-Fit (RMFF): (Dhall and Liu, 1978)

- Let the processors be indexed as  $\mu_1, \mu_2, \dots, \mu_m$
- Assign tasks in order of increasing periods (i.e., RM order).
- For each task  $\tau_i$ , choose the lowest previously-used  $j$  such that  $\tau_i$ , together with all tasks that have already been assigned to processor  $\mu_j$ , can be feasibly scheduled according to the utilization-based RM-feasibility test.

If all tasks are successfully assigned using RMFF, then the tasks are schedulable on  $m$  processors.

# Partitioned scheduling

Processor utilization analysis for RMFF:

- A sufficient condition for partitioned RMFF scheduling of synchronous task sets with  $n$  tasks on  $m$  processors is

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq m(2^{1/2} - 1)$$

(Oh & Baker, 1998)

Note:  $U_{RMFF} = m(2^{1/2} - 1) \approx 0.41m$

Thus: task sets whose utilization do not exceed  $\approx 41\%$  of the total processor capacity is always RMFF-schedulable.

# Partitioned scheduling

Processor utilization analysis for RMFF:

1. All tasks are independent
2. All tasks are periodic or sporadic
3. All tasks have identical offsets (= synchronous task set)
4. Task deadline **equals** the period (= implicit-deadline tasks)
5. Task preemptions are allowed
6. All processors are **identical**
7. Task migrations are not allowed

# Global scheduling

## General characteristics:

- All ready tasks are kept in a common (global) queue that is shared among the processors
- Whenever a processor becomes idle, a task from the global queue is selected for execution on that processor.
- After being preempted, a task may be dispatched to a processor other than the one that started executing the task.

# Global scheduling

## Advantages:

- Supported by most multiprocessor operating systems
  - Windows, macOS, Linux, ...
- Effective utilization of processing resources
  - Unused processor time can easily be reclaimed, for example when a task does not execute its full WCET.

## Disadvantages:

- Weak theoretical framework
  - Few results from the single-processor analysis can be used

# Weak theoretical framework

The "root of all evil" in global scheduling: (Liu, 1969)

Few of the results obtained for a single processor generalize directly to the multiple processor case; bringing in additional processors adds a new dimension to the scheduling problem. The simple fact that *a task can use only one processor even when several processors are free at the same time* adds a surprising amount of difficulty to the scheduling of multiple processors.

# Weak theoretical framework

## Underlying causes:

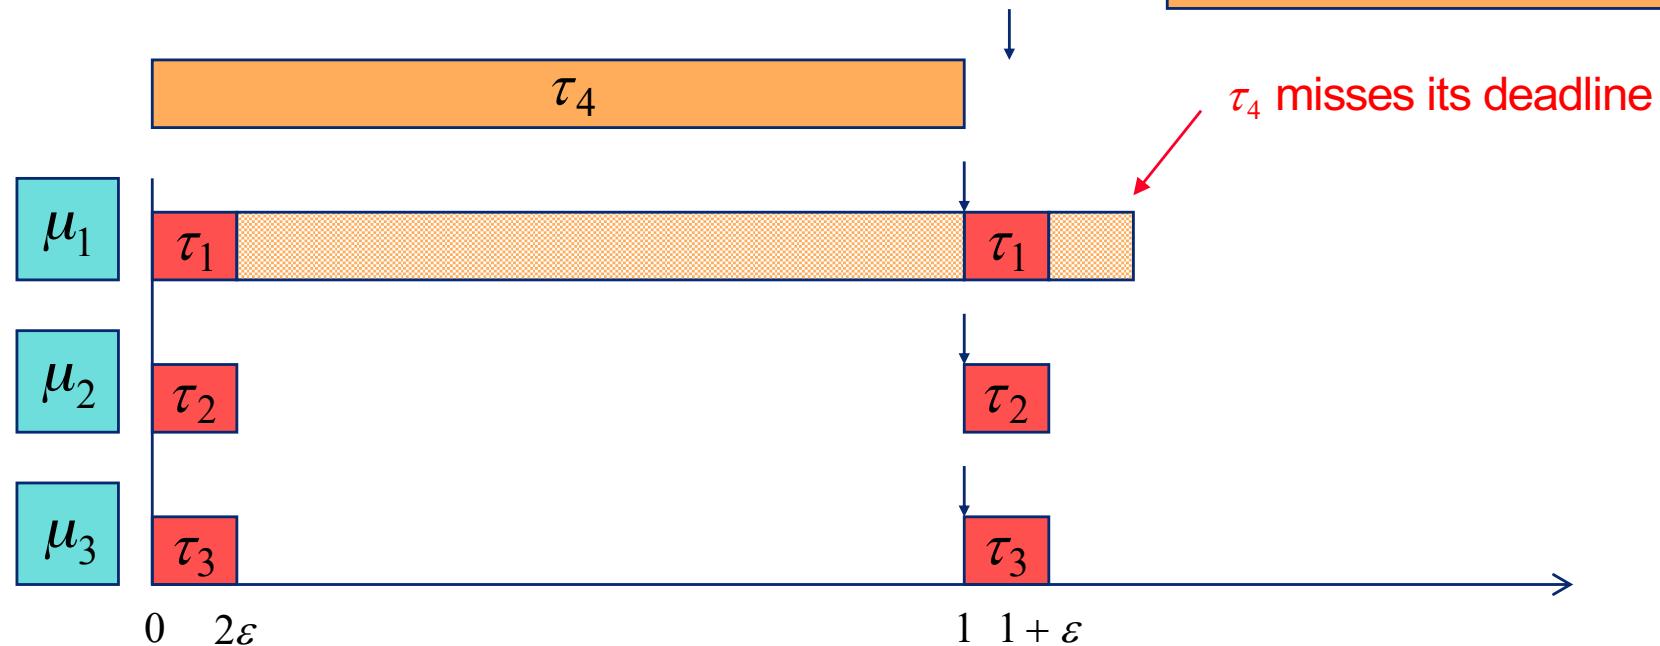
- Dhall's effect:
  - With RM, DM and EDF, some low-utilization task sets can be non-schedulable regardless of how many processors are used. Thus, any utilization guarantee bound would become so low that it would be useless in practice.
  - This is in contrast to the single-processor case, where we have utilization guarantee bounds of 69.3% (RM) and 100% (EDF).
- Hard-to-find critical instant:
  - A critical instant does not always occur when a task arrives at the same time as all its higher-priority tasks.
  - This is in contrast to the single-processor case with RM and DM (and any other static-priority policy).

# Weak theoretical framework

Dhall's effect: (Dhall & Liu, 1978)

(RM scheduling)

$$\begin{aligned}\tau_1 : & \{ C_1 = 2\epsilon, T_1 = 1 \} \\ \tau_2 : & \{ C_2 = 2\epsilon, T_2 = 1 \} \\ \tau_3 : & \{ C_3 = 2\epsilon, T_3 = 1 \} \\ \tau_4 : & \{ C_4 = 1, T_4 = 1 + \epsilon \}\end{aligned}$$



# Weak theoretical framework

## Dhall's effect:

- Applies for RM, DM and EDF scheduling
- The utilization of a non-schedulable task set can be as low as to 1 (= 100%) no matter how many processors are used.

$$U_{global} = m \frac{2\epsilon}{1} + \frac{1}{1+\epsilon} \rightarrow 1$$

when  $\epsilon \rightarrow 0$

Note: Total available processor capacity is  $m$  (=  $m \cdot 100\%$ )

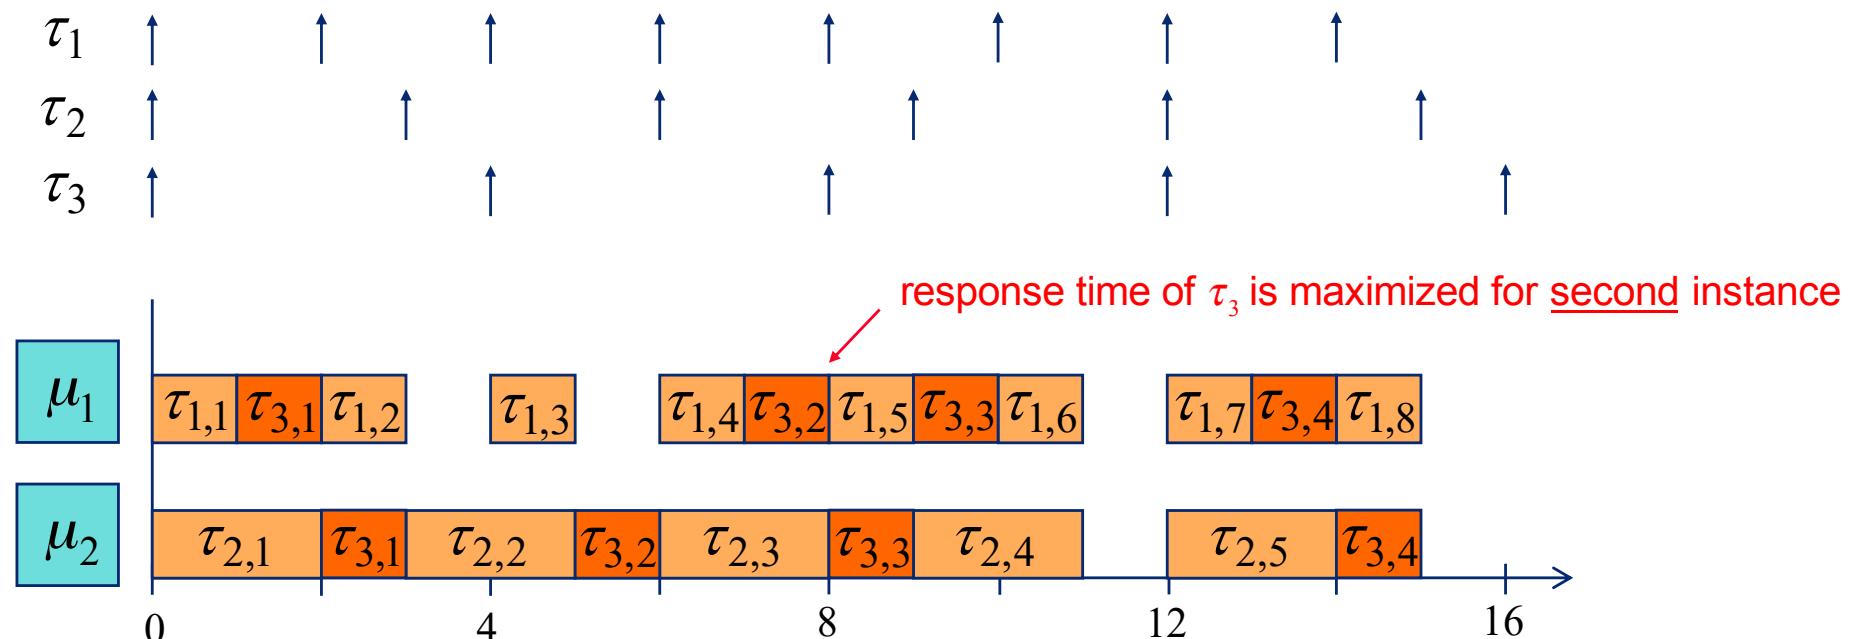
## Consequence:

New multiprocessor priority-assignment schemes are needed!

# Weak theoretical framework

Hard-to-find critical instant:  
(RM scheduling)

$$\begin{aligned}\tau_1 &: \{ C_1 = 1, T_1 = 2 \} \\ \tau_2 &: \{ C_2 = 2, T_2 = 3 \} \\ \tau_3 &: \{ C_3 = 2, T_3 = 4 \}\end{aligned}$$



# Weak theoretical framework

## Hard-to-find critical instant:

- A critical instant does not always occur when a task arrives at the same time as all its higher-priority tasks.
- Finding the critical instant is, in general, a problem with exponential time complexity
- Note: recall that knowledge about an easy-to-find critical instant is a fundamental assumption in the single-processor feasibility tests for static-priority scheduling.

## Consequence:

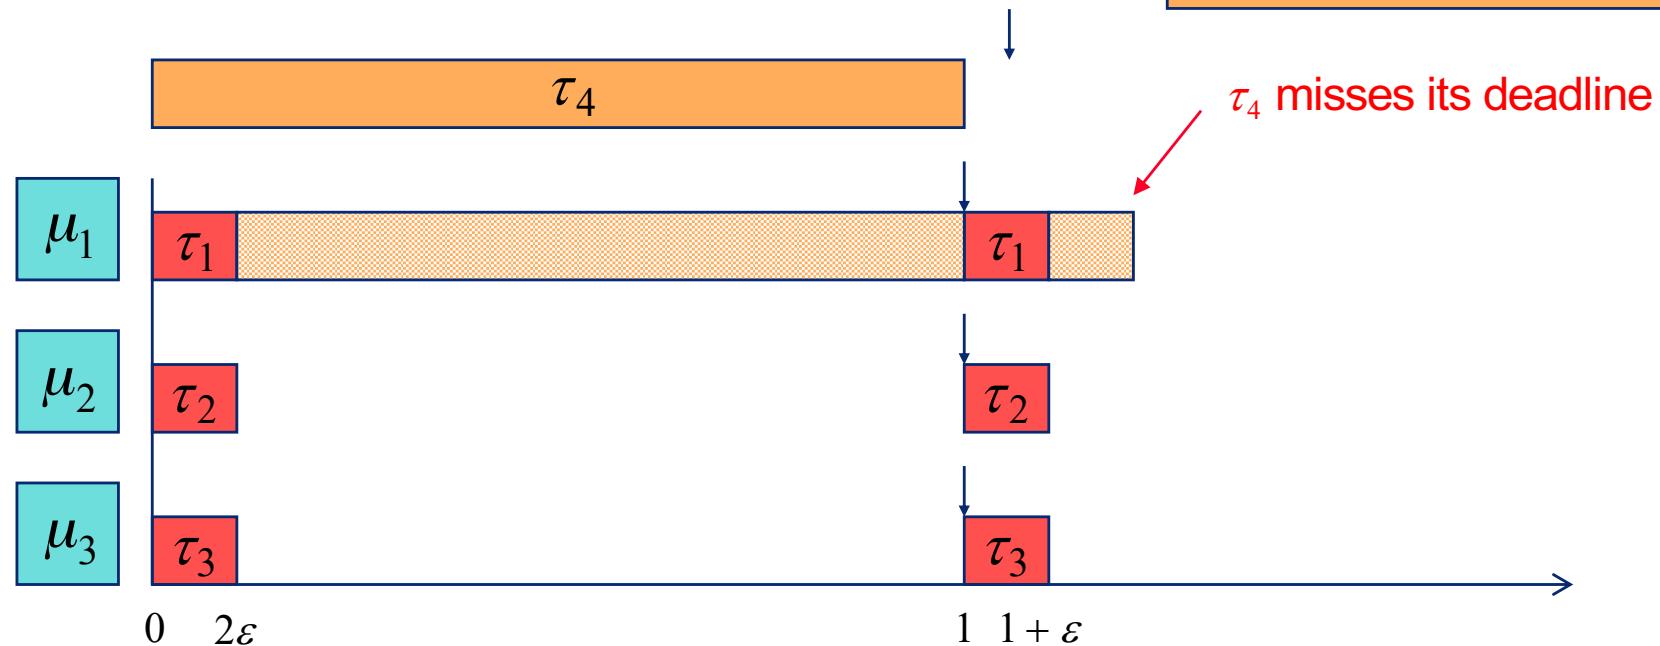
New methods for constructing effective multiprocessor feasibility tests are needed!

# Weak theoretical framework

Dhall's effect: (Dhall & Liu, 1978)

(RM scheduling)

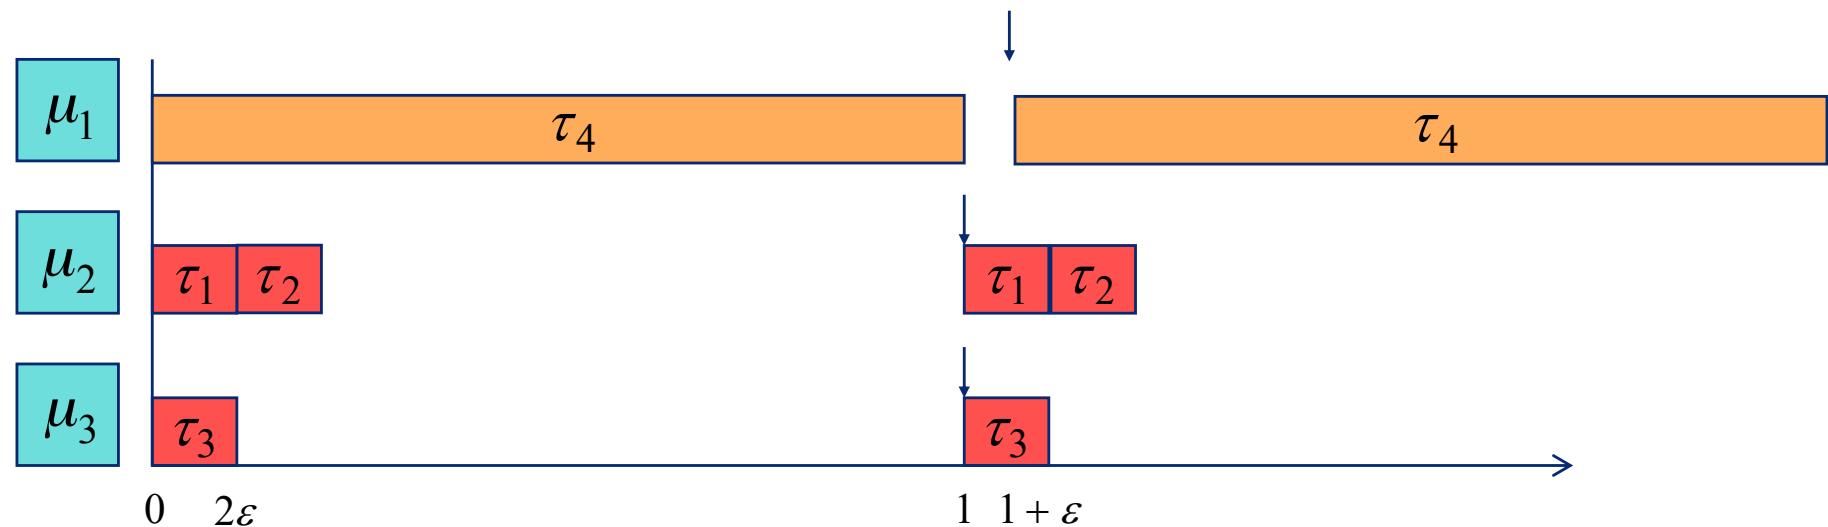
$$\begin{aligned}\tau_1 : & \{ C_1 = 2\epsilon, T_1 = 1 \} \\ \tau_2 : & \{ C_2 = 2\epsilon, T_2 = 1 \} \\ \tau_3 : & \{ C_3 = 2\epsilon, T_3 = 1 \} \\ \tau_4 : & \{ C_4 = 1, T_4 = 1 + \epsilon \}\end{aligned}$$



# New priority-assignment scheme

How to avoid Dhall's effect:

- Problem: RM, DM & EDF only account for task deadlines!  
Actual computation demands are not accounted for.
- Solution: Dhall's effect can easily be avoided by letting tasks with high utilization receive higher priority:



# New priority-assignment scheme

RM-US[m/(3m-2)]: (Andersson, Baruah & Jonsson, 2001)

- RM-US[m/(3m-2)] assigns (static) priorities to tasks according to the following rule:

If  $U_i > m/(3m - 2)$  then  $\tau_i$  has the highest priority  
(ties broken arbitrarily)

If  $U_i \leq m/(3m - 2)$  then  $\tau_i$  has RM priority

- Clearly, tasks with higher utilization  $U_i = C_i / T_i$  get higher priority.

# Example: RM-US[m/(3m-2)]

## RM-US[m/(3m-2)] example:

Assign priorities according to RM-US[m/(3m-2)], assuming the following task set to be scheduled on 3 processors:

$$\tau_1 : \{ C_1 = 1, T_1 = 7 \} \quad \tau_2 : \{ C_2 = 2, T_2 = 10 \}$$

$$\tau_3 : \{ C_3 = 9, T_3 = 20 \} \quad \tau_4 : \{ C_4 = 11, T_4 = 22 \}$$

$$\tau_5 : \{ C_5 = 2, T_5 = 25 \}$$

# New feasibility tests

Processor utilization analysis for RM-US[m/(3m-2)]:

- A sufficient condition for RM-US[m/(3m-2)] scheduling of synchronous task sets with  $n$  tasks on  $m$  processors is

$$U = \sum_{i=1}^n \frac{C_i}{T_i} \leq \frac{m^2}{3m-2}$$

(Andersson, Baruah  
& Jonsson, 2001)

Question: does RM-US[m/(3m-2)] avoid Dhall's effect?

# New feasibility tests

## Processor utilization analysis for RM-US[m/(3m-2)]:

- We observe that, regardless of the number of processors, the task set will always meet its deadlines as long as no more than one third of the processing capacity is used:

$$U_{RM-US[m/(3m-2)]} = \lim_{m \rightarrow \infty} \frac{m^2}{3m-2} = \frac{m}{3}$$

- RM-US[m/(3m-2)] thus avoids Dhall's effect since we can always add more processors if deadlines were missed.
- Note that this remedy was not possible with traditional RM.

# New feasibility tests

Processor utilization analysis for RM-US[m/(3m-2)]:

1. All tasks are independent
2. All tasks are **periodic** (i.e. not applicable for sporadic tasks)
3. All tasks have identical offsets (= synchronous task set)
4. Task deadline **equals** the period (= implicit-deadline tasks)
5. Task preemptions are allowed
6. All processors are **identical**
7. Task migrations are **allowed**

# New feasibility tests

## Early breakthrough results in global scheduling:

- Static priorities:
  - 2001: RM-US $[m/(3m-2)]$  is proven to circumvent Dhall's effect and to have a non-zero guarantee bound of  $m/(3m-2) \geq 33.3\%$ .
  - 2008: Andersson proposed the SM-US $\{2/(3+\sqrt{5})\}$  scheme with a guarantee bound of  $2/(3+\sqrt{5}) = 38.2\%$ .
- Dynamic priorities:
  - 2002: Srinivasan & Baruah proposed the EDF-US $[m/(2m-1)]$  scheme with a guarantee bound of  $m/(2m-1) \geq 50\%$ .
- Optimal multiprocessor scheduling:
  - 1996: Baruah *et al.* proposed p-fair (priorities + time quanta) scheduling and dynamic task priorities as an approach to achieve a guarantee bound of 100% on a multiprocessor.

# New feasibility tests

## Response-time analysis for multiprocessors:

- Uses the same principle as the single-processor case, where the response time for a task  $\tau_i$  consists of:

$C_i$  The task's uninterrupted execution time (WCET)

$I_i$  Interference from higher-priority tasks

$$R_i = C_i + I_i$$

- The difference is that the calculation of interference now has to account for the fact that higher-priority tasks can execute in parallel on the processors.

# New feasibility tests

## Response-time analysis for multiprocessors:

- For the multiprocessor case, with  $n$  tasks and  $m$  processors, we observe two things:
  1. Interference can only occur when  $n > m$ .
  2. Interference can only affect the  $n - m$  tasks with lowest priority since the  $m$  highest-priority tasks will always execute in parallel without contention on the  $m$  processors.
- Consequently, interference of a task is a function of the execution overlap of its higher-priority tasks.

# New feasibility tests

## Response-time analysis for multiprocessors:

- The following two observations give us the secret to analyzing the interference of a task:

With respect to the execution overlap it can be shown that the interference is maximized when the higher-priority tasks completely overlap their execution.

Compared to the single-processor case, one extra instance of each higher-priority task must be accounted for in the interference analysis.

(due to the uncertainty regarding the critical instant).

# New feasibility tests

## Response-time analysis for multiprocessors:

- The worst-case interference term is

$$I_i = \frac{1}{m} \sum_{\forall j \in hp(i)} \left( \left\lceil \frac{R_i}{T_j} \right\rceil \cdot C_j + C_j \right)$$

where  $hp(i)$  is the set of tasks with higher priority than  $\tau_i$ .

- The worst-case response time for a task  $\tau_i$  is thus:

$$R_i = C_i + \frac{1}{m} \sum_{\forall j \in hp(i)} \left( \left\lceil \frac{R_i}{T_j} \right\rceil \cdot C_j + C_j \right)$$

# New feasibility tests

## Response-time analysis for multiprocessors:

- As before, an iterative approach can be used for finding the worst-case response time:

$$R_i^{n+1} = C_i + \frac{1}{m} \sum_{\forall j \in hp(i)} \left( \left\lceil \frac{R_i^n}{T_j} \right\rceil \cdot C_j + C_j \right)$$

- We now have a sufficient condition for static-priority scheduling of periodic tasks on identical processors:

$$\forall i: R_i \leq D_i$$

# New feasibility tests

Response-time analysis for multiprocessors:

1. All tasks are independent
2. All tasks are periodic (i.e. not applicable for sporadic tasks)
3. All tasks have identical offsets (= synchronous task set)
4. Task deadline does not exceed the period  
(= constrained-deadline tasks)
5. Task preemptions are allowed
6. All processors are identical
7. Task migrations are allowed